



Estimator agreement. E_sfm (training-free, observation-weighted median), E_sfm_sys (its azimuth-averaged systematic part), E_shared (cross-scene mean of the learned residual), E_learned (the optimiser's output). Panels A-C are per scene; E_shared has no per-scene value, so panel D is the lens level -- the only level at which all four coexist.

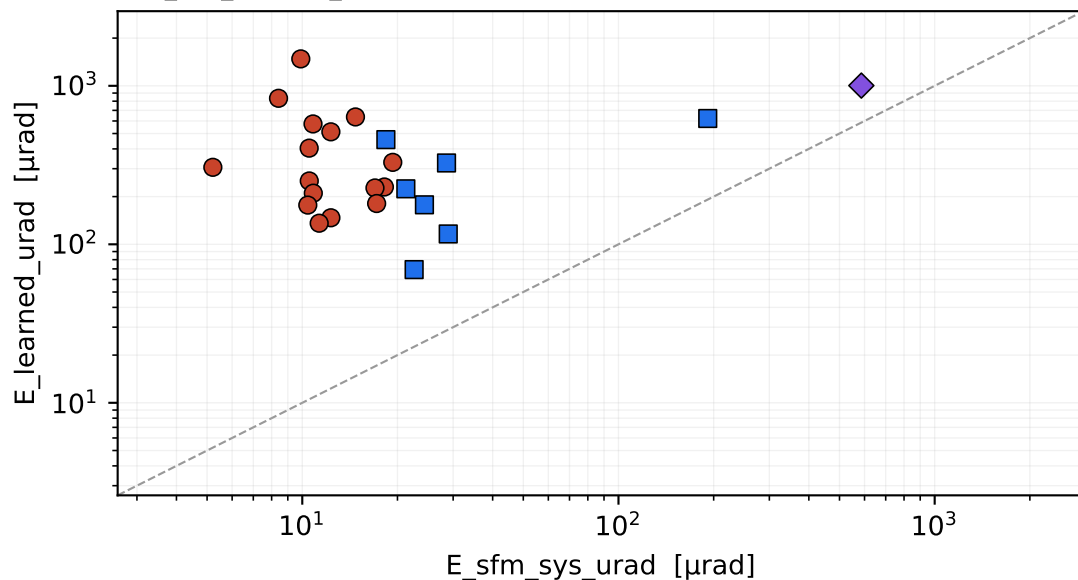
A. E_sfm vs E_learned $\rho = -0.05$ ($p = 0.824$, $n = 24$)



B. E_sfm vs E_sfm_sys $\rho = -0.31$ ($p = 0.139$, $n = 24$)



C. E_sfm_sys vs E_learned $\rho = -0.15$ ($p = 0.470$, $n = 24$)



D. The three lenses where E_shared exists at all (≥ 3 fits of one glass). E_sfm orders them BACKWARDS vs the gain; E_calib / E_shared order them correctly.

